

ROV Vision System — Degradation Analysis Report

1 How Confidence Scores Changed Under Different Conditions

Condition	Prediction	Confidence	Latency (ms)	Notes
Baseline (Clean)	coral reef / fish	~82%	~18 ms	High confidence; correct category
Turbidity	hen-of-the-woods mushroom	~34%	~18 ms	Blur destroyed spatial features; wrong class
Color Shift	submarine	~28%	~18 ms	Blue cast shifts prediction domain entirely
Sensor Noise	coral reef	~41%	~18 ms	Class retained but confidence drops ~50%

Under the clean baseline the model assigned high confidence (~82%) to the correct aquatic category. Confidence dropped substantially under every simulated degradation:

- Turbidity (–58 pp): The Gaussian blur + haze overlay eliminated the fine-grained texture cues ResNet18 relies on for aquatic recognition, causing a misclassification into an entirely different visual domain (fungi).
- Color Shift (–54 pp): Removing the red channel and boosting blue radically changed the pixel distribution, pushing activation toward non-aquatic, blue-dominant classes (e.g. submarine, tench).
- Sensor Noise (–41 pp): Gaussian noise degraded edge and texture information, halving confidence while occasionally preserving the correct top-1 class, revealing that noise is less catastrophic than colour/structure degradation.

2 Why the Model Failed on Specific Degradations

2.1 Turbidity

ResNet18 learnt to classify objects through texture gradients and spatial frequency patterns. A 25×25 Gaussian blur collapses these high-frequency components, leaving only low-frequency blobs. The resulting feature map looks nothing like the training distribution, so the model latches onto the nearest low-frequency pattern in ImageNet — often fungal caps or abstract textures.

2.2 Depth Color Shift

ImageNet images are RGB-balanced. Attenuating the red channel to 15% and saturating blue creates an out-of-distribution colour histogram. Early convolutional layers extract colour blobs, and the resulting blue-dominant activations match ship hulls or lake-surface images rather than fish schools. This is a domain-shift failure.

2.3 Sensor Noise

Gaussian noise ($\sigma = 35$) corrupts pixel values by ± 35 DN on average. While edges are partially preserved, fine texture features become unreliable. The model maintains the correct broad class more often than in the other two conditions, confirming that structure matters more than raw pixel values — but confidence still falls sharply because the neural network was never trained on noisy imagery.

3 Is ResNet18 Suitable for Real-Time Edge Hardware?

At ~18 ms per inference on a standard CPU, ResNet18 just meets the 33.33 ms budget required for 30 FPS. However, this measurement was taken on a desktop-class CPU with no concurrency. On edge platforms typically deployed in ROVs (e.g. NVIDIA Jetson Nano, Raspberry Pi 4) inference time rises to 60–200 ms without hardware optimisation, which would cause systematic frame drops.

Additionally, ResNet18 was not trained on underwater imagery and shows severe confidence collapse under all three ocean-relevant degradations. Therefore:

- ResNet18 is NOT suitable as-is for real-time edge deployment in underwater ROVs.

- For latency alone it can work on GPU-equipped edge boards (Jetson AGX ~5 ms) but fails on CPU-only embedded platforms.
- Model quality (accuracy under degradation) is the primary disqualifier, independent of latency.

4 Engineering Steps to Make the System More Robust

Step 1 — Domain-Adaptive Fine-Tuning on Underwater Data

Collect or synthesise a labelled underwater dataset (FathomNet, MBARI DeepSea, or synthetic rendering) covering turbidity, colour-shifted, and noisy conditions. Fine-tune the final 2–3 ResNet18 blocks on this data while freezing early layers. This directly closes the distribution gap responsible for Steps 2 and 3 failures. Expected confidence improvement: +30–50 percentage points under simulated conditions.

Step 2 — Image Pre-Processing Pipeline (Underwater Enhancement)

Insert a real-time pre-processing stage before model inference:

1. White-balance correction using an adaptive grey-world algorithm to restore red-channel content lost to depth absorption.
2. Contrast-Limited Adaptive Histogram Equalization (CLAHE) to counteract turbidity-induced low contrast.
3. Bilateral or guided-filter denoising (preserves edges, unlike Gaussian) to suppress sensor noise without destroying texture.

These transforms bring degraded frames closer to the training distribution at near-zero compute cost (~1–2 ms on CPU).

5 Bonus — MobileNetV2 vs ResNet18

MobileNetV2 uses depthwise-separable convolutions to achieve ~8× fewer multiply-add operations than ResNet18. Expected inference on the same CPU: 6–10 ms vs 18 ms, comfortably within the 33 ms / 30 FPS budget. However, MobileNetV2 also fails under domain shift unless fine-tuned. The recommended architecture for a production ROV system is MobileNetV3-Small fine-tuned on underwater data, TensorRT-quantised to INT8, and deployed on a Jetson Orin Nano targeting < 5 ms per frame with hardware-pipelined camera capture to eliminate I/O bottlenecks.

Frame-Drop Analysis

At 30 FPS, each frame must be processed within 33.33 ms. ResNet18 at ~18 ms leaves ~15 ms for I/O and post-processing — marginal on CPU, safe on GPU. If latency exceeds the budget (common on embedded CPUs without GPU), the pipeline must either: (a) drop to 15 FPS, (b) use asynchronous inference with a frame queue, or (c) switch to a lighter model. Hardware suggestions: NVIDIA Jetson Orin (DLA accelerator, ~2 ms/frame) or Google Coral Edge TPU (INT8, ~3 ms/frame) for true real-time performance.

— End of Report —